


Aegis — Governed-Agent Platform • Deployment Architecture

One control plane, every agent through it • deny-by-default • Built on AWS • customer-safe reference architecture (see NOT-CLAIMS)

Deploys with one command (cdk deploy)
At rest = \$1-3/month • August 2026

CUSTOMER AWS ACCOUNT — your region • your keys • your IdP • GovCloud variant included



Aegis — A Governed Request, Step by Step

Three journeys through the same gateway: the routine request, the consequential action, and the bad day

Every step below writes an audit row
Mock flow · illustrative timings & costs

A • Routine request — a caseworker screens an applicant (~3 seconds · ≈ \$0.007)

1 Sign-in Caseworker signs in via the agency IdP → Cognito issues a JWT with role <i>caseworker</i>.	2 Ask "Screen applicant #4141 for SNAP eligibility." The agent plans a call to its declared tool <i>screen_applicant</i>.	3 Gateway checks Kill switch clear ✓ · JWT verified ✓ · tool declared in manifest ✓ · agent n caseworker permits ✓ · budget preflight ✓	4 Masking Applicant record crosses the boundary — SSN and phone are anonymized <i>before</i> anything reaches the model. Fail-closed: if masking can't run, the call denies.	5 Grounded inference Guardrail applied on input; model answers against the pinned FPL reference data; guardrail applied again on output.	6 Answer + evidence "Likely eligible — household of 3, income 118% FPL" with the cited basis. Audit row written; tokens metered to the department.
---	--	---	---	---	---

B • Consequential action — submitting the determination (the human gate)

1 Agent proposes <i>submit_determination</i> is NOT in the agent's own tool grants — it exists only behind the gate. The AI cannot do this alone, by construction.	2 Workflow pauses Step Functions parks the case on a task token. Nothing is submitted. SNS notifies the approver role.	3 Second human reviews A supervisor — not the requester — sees the case, the evidence, and the masked record, and approves.	4 Single-use approval The approval is bound to this agent + this tool + this case, expires in 1 hour, and is consumed exactly once.	5 Action executes Determination submitted. Audit row carries both humans, the approval id, and the timestamp chain.	6 The negatives hold Self-approval → DENY. Replayed approval → DENY. Expired token → DENY. Each denial is itself audited.
---	---	---	--	--	--

C • The bad day — deny-by-default and the kill switch

1 Undeclared tool A modified agent tries <i>db.export_all</i> — not in its manifest. DENY before any data moves; the attempt is in the audit.	2 Budget breach A looping agent hits its hard token cap. Next call: DENY — a denied call, not a surprise invoice. Alarm fired at 85%.	3 2:00 AM — incident Operator engages the Kill Switch: one parameter write. Every agent action platform-wide now denies with reason KILL_SWITCH_ENGAGED.	4 Containment proven Canary confirms the deny. The engage event, actor, and reason are already in the audit trail.	5 Release — different human Separation of duties: the identity that engaged cannot release. A second operator disengages; service resumes.	6 The morning after The full incident — engage, every denial, release — is a query, not a forensic project. Evidence filed to WORM.
---	--	--	---	---	--

The line that closes the meeting: every arrow on this page is a row in an append-only, hash-chained audit table that nobody — including administrators — can edit or delete. "Can you prove it?" is answered with a query.